





```
Train set length: 4072
Test set length: 174
Model: "model"

Layer (Type) Output Shape Param # Connected to
-----
Input_1 (InputLayer) (None, 131) 0 input_1[0][0]
Input_2 (InputLayer) (None, 6) 0
concatenate (Concatenate) (None, 7) 0 dense_1[0][0]
batch_normalization (BatchNorm) (None, 7) 28 concatenate[0][0]
dropout_1 (Dropout) (None, 7) 0 batch_normalization[0][0]
dense_1 (Dense) (None, 15) 120 dropout_1[0][0]
dropout_2 (Dropout) (None, 15) 0 dense_1[0][0]
dense_2 (Dense) (None, 15) 240 dropout_1[0][0]
dropout_2 (Dropout) (None, 15) 0 dense_2[0][0]
dense_3 (Dense) (None, 1) 16 dropout_2[0][0]
-----
Total params: 418
Trainable params: 404
Non-trainable params: 14

Fold 1 *****
Epoch 1/50 ===== - 2s 1ms/step - loss: 0.9346 - root_mean_squared_error: 0.9668 - val_loss: 0.6546 - val_root_mean_squared_error: 0.8090
Epoch 0001: val_root_mean_squared_error improved from inf to 0.80904, saving model to model_0.h5
Epoch 2/50 ===== - 3s 2ms/step - loss: 0.7764 - root_mean_squared_error: 0.8812 - val_loss: 0.4947 - val_root_mean_squared_error: 0.7033
Epoch 0002: val_root_mean_squared_error improved from 0.80904 to 0.70334, saving model to model_0.h5
Epoch 3/50 ===== - 3s 2ms/step - loss: 0.7190 - root_mean_squared_error: 0.8480 - val_loss: 0.3404 - val_root_mean_squared_error: 0.5835
Epoch 0003: val_root_mean_squared_error improved from 0.70334 to 0.58346, saving model to model_0.h5
Epoch 4/50 ===== - 2s 1ms/step - loss: 0.6162 - root_mean_squared_error: 0.7850 - val_loss: 0.2687 - val_root_mean_squared_error: 0.5164
Epoch 0004: val_root_mean_squared_error improved from 0.58346 to 0.51641, saving model to model_0.h5
Epoch 5/50 ===== - 2s 1ms/step - loss: 0.5049 - root_mean_squared_error: 0.7106 - val_loss: 0.1694 - val_root_mean_squared_error: 0.4576
Epoch 0005: val_root_mean_squared_error improved from 0.51641 to 0.45757, saving model to model_0.h5
Epoch 6/50 ===== - 2s 1ms/step - loss: 0.4800 - root_mean_squared_error: 0.6922 - val_loss: 0.1484 - val_root_mean_squared_error: 0.3853
Epoch 0006: val_root_mean_squared_error improved from 0.45757 to 0.38529, saving model to model_0.h5
Epoch 7/50 ===== - 2s 1ms/step - loss: 0.4207 - root_mean_squared_error: 0.6486 - val_loss: 0.1230 - val_root_mean_squared_error: 0.3507 - root_mean_squared_error: 0.4413
Epoch 0007: val_root_mean_squared_error improved from 0.38529 to 0.35070, saving model to model_0.h5
Epoch 8/50 ===== - 2s 1ms/step - loss: 0.3986 - root_mean_squared_error: 0.6313 - val_loss: 0.1000 - val_root_mean_squared_error: 0.3163
Epoch 0008: val_root_mean_squared_error improved from 0.35070 to 0.31627, saving model to model_0.h5
Epoch 9/50 ===== - 2s 1ms/step - loss: 0.3532 - root_mean_squared_error: 0.5918 - val_loss: 0.0788 - val_root_mean_squared_error: 0.2808
Epoch 0009: val_root_mean_squared_error improved from 0.31627 to 0.28075, saving model to model_0.h5
Epoch 10/50 ===== - 2s 1ms/step - loss: 0.2931 - root_mean_squared_error: 0.5414 - val_loss: 0.0281 - val_root_mean_squared_error: 0.2556
Epoch 0010: val_root_mean_squared_error improved from 0.28075 to 0.25557, saving model to model_0.h5
Epoch 11/50 ===== - 2s 1ms/step - loss: 0.2914 - root_mean_squared_error: 0.5398 - val_loss: 0.0201 - val_root_mean_squared_error: 0.2283
Epoch 0011: val_root_mean_squared_error improved from 0.25557 to 0.22829, saving model to model_0.h5
Epoch 12/50 ===== - 2s 1ms/step - loss: 0.2482 - root_mean_squared_error: 0.4982 - val_loss: 0.0403 - val_root_mean_squared_error: 0.2089
Epoch 0012: val_root_mean_squared_error improved from 0.22829 to 0.20080, saving model to model_0.h5
Epoch 13/50 ===== - 2s 1ms/step - loss: 0.2354 - root_mean_squared_error: 0.4852 - val_loss: 0.0331 - val_root_mean_squared_error: 0.1818
Epoch 0013: val_root_mean_squared_error improved from 0.20080 to 0.18182, saving model to model_0.h5
Epoch 14/50 ===== - 2s 1ms/step - loss: 0.2248 - root_mean_squared_error: 0.4742 - val_loss: 0.0281 - val_root_mean_squared_error: 0.1675
Epoch 0014: val_root_mean_squared_error improved from 0.18182 to 0.16752, saving model to model_0.h5
Epoch 15/50 ===== - 2s 1ms/step - loss: 0.2033 - root_mean_squared_error: 0.4509 - val_loss: 0.0214 - val_root_mean_squared_error: 0.1463
Epoch 0015: val_root_mean_squared_error improved from 0.16752 to 0.14628, saving model to model_0.h5
Epoch 16/50 ===== - 2s 1ms/step - loss: 0.1848 - root_mean_squared_error: 0.4299 - val_loss: 0.0164 - val_root_mean_squared_error: 0.1280
Epoch 0016: val_root_mean_squared_error improved from 0.14628 to 0.12805, saving model to model_0.h5
Epoch 17/50 ===== - 2s 1ms/step - loss: 0.1638 - root_mean_squared_error: 0.4047 - val_loss: 0.0147 - val_root_mean_squared_error: 0.1211
Epoch 0017: val_root_mean_squared_error improved from 0.12805 to 0.12114, saving model to model_0.h5
Epoch 18/50 ===== - 2s 1ms/step - loss: 0.1540 - root_mean_squared_error: 0.3924 - val_loss: 0.0211 - val_root_mean_squared_error: 0.1099
Epoch 0018: val_root_mean_squared_error improved from 0.12114 to 0.10987, saving model to model_0.h5
Epoch 19/50 ===== - 2s 1ms/step - loss: 0.1461 - root_mean_squared_error: 0.3822 - val_loss: 0.0105 - val_root_mean_squared_error: 0.1025
Epoch 0019: val_root_mean_squared_error improved from 0.10987 to 0.10254, saving model to model_0.h5
Epoch 20/50 ===== - 2s 1ms/step - loss: 0.1249 - root_mean_squared_error: 0.3535 - val_loss: 0.0083 - val_root_mean_squared_error: 0.0939
Epoch 0020: val_root_mean_squared_error improved from 0.10254 to 0.09094, saving model to model_0.h5
Epoch 21/50 ===== - 2s 1ms/step - loss: 0.1065 - root_mean_squared_error: 0.3441 - val_loss: 0.0017 - val_root_mean_squared_error: 0.0887
Epoch 0021: val_root_mean_squared_error improved from 0.09094 to 0.08770, saving model to model_0.h5
Epoch 22/50 ===== - 2s 1ms/step - loss: 0.1065 - root_mean_squared_error: 0.3264 - val_loss: 0.0059 - val_root_mean_squared_error: 0.0770
Epoch 0022: val_root_mean_squared_error improved from 0.08770 to 0.07695, saving model to model_0.h5
Epoch 23/50 ===== - 2s 1ms/step - loss: 0.0993 - root_mean_squared_error: 0.3151 - val_loss: 0.0049 - val_root_mean_squared_error: 0.0702
Epoch 0023: val_root_mean_squared_error improved from 0.07695 to 0.07021, saving model to model_0.h5
Epoch 24/50 ===== - 2s 1ms/step - loss: 0.0954 - root_mean_squared_error: 0.3088 - val_loss: 0.0044 - val_root_mean_squared_error: 0.0663
Epoch 0024: val_root_mean_squared_error improved from 0.07021 to 0.06628, saving model to model_0.h5
Epoch 25/50 ===== - 2s 1ms/step - loss: 0.0874 - root_mean_squared_error: 0.2956 - val_loss: 0.0035 - val_root_mean_squared_error: 0.0613
Epoch 0025: val_root_mean_squared_error improved from 0.06628 to 0.06134, saving model to model_0.h5
Epoch 26/50 ===== - 2s 1ms/step - loss: 0.0821 - root_mean_squared_error: 0.2866 - val_loss: 0.0034 - val_root_mean_squared_error: 0.0582
Epoch 0026: val_root_mean_squared_error improved from 0.06134 to 0.05824, saving model to model_0.h5
Epoch 27/50 ===== - 2s 1ms/step - loss: 0.0694 - root_mean_squared_error: 0.2815 - val_loss: 0.0026 - val_root_mean_squared_error: 0.0551
Epoch 0027: val_root_mean_squared_error improved from 0.05824 to 0.05508, saving model to model_0.h5
Epoch 28/50 ===== - 2s 1ms/step - loss: 0.0694 - root_mean_squared_error: 0.2634 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0510
Epoch 0028: val_root_mean_squared_error improved from 0.05508 to 0.05442, saving model to model_0.h5
Epoch 29/50 ===== - 2s 1ms/step - loss: 0.0653 - root_mean_squared_error: 0.2555 - val_loss: 0.0028 - val_root_mean_squared_error: 0.0480
Epoch 0029: val_root_mean_squared_error improved from 0.05442 to 0.04999, saving model to model_0.h5
Epoch 30/50 ===== - 2s 1ms/step - loss: 0.0598 - root_mean_squared_error: 0.2446 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0463
Epoch 0030: val_root_mean_squared_error improved from 0.04999 to 0.04628, saving model to model_0.h5
Epoch 31/50 ===== - 3s 2ms/step - loss: 0.0566 - root_mean_squared_error: 0.2380 - val_loss: 0.0021 - val_root_mean_squared_error: 0.0457
Epoch 0031: val_root_mean_squared_error improved from 0.04628 to 0.04566, saving model to model_0.h5
Epoch 32/50 ===== - 3s 2ms/step - loss: 0.0477 - root_mean_squared_error: 0.2185 - val_loss: 0.0020 - val_root_mean_squared_error: 0.0450
Epoch 0032: val_root_mean_squared_error improved from 0.04566 to 0.04497, saving model to model_0.h5
Epoch 33/50 ===== - 3s 2ms/step - loss: 0.0505 - root_mean_squared_error: 0.2247 - val_loss: 0.0010 - val_root_mean_squared_error: 0.0448
Epoch 0033: val_root_mean_squared_error improved from 0.04497 to 0.04483, saving model to model_0.h5
Epoch 34/50 ===== - 3s 2ms/step - loss: 0.0478 - root_mean_squared_error: 0.2187 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0426
Epoch 0034: val_root_mean_squared_error improved from 0.04483 to 0.04423, saving model to model_0.h5
Epoch 35/50 ===== - 3s 2ms/step - loss: 0.0424 - root_mean_squared_error: 0.2058 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0426
Epoch 0035: val_root_mean_squared_error did not improve from 0.04251
Epoch 36/50 ===== - 2s 1ms/step - loss: 0.0404 - root_mean_squared_error: 0.2010 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0416
Epoch 0036: val_root_mean_squared_error did not improve from 0.04251
Epoch 37/50 ===== - 2s 1ms/step - loss: 0.0364 - root_mean_squared_error: 0.1908 - val_loss: 0.0017 - val_root_mean_squared_error: 0.0416
Epoch 0037: val_root_mean_squared_error improved from 0.04251 to 0.04162, saving model to model_0.h5
Epoch 38/50 ===== - 2s 1ms/step - loss: 0.0343 - root_mean_squared_error: 0.1853 - val_loss: 0.0017 - val_root_mean_squared_error: 0.0414
Epoch 0038: val_root_mean_squared_error improved from 0.04162 to 0.04142, saving model to model_0.h5
Epoch 39/50 ===== - 2s 1ms/step - loss: 0.0326 - root_mean_squared_error: 0.1806 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0410
Epoch 0039: val_root_mean_squared_error improved from 0.04142 to 0.04095, saving model to model_0.h5
Epoch 40/50 ===== - 2s 1ms/step - loss: 0.0311 - root_mean_squared_error: 0.1763 - val_loss: 0.0016 - val_root_mean_squared_error: 0.0396
Epoch 0040: val_root_mean_squared_error improved from 0.04095 to 0.03956, saving model to model_0.h5
Epoch 41/50 ===== - 2s 1ms/step - loss: 0.0279 - root_mean_squared_error: 0.1672 - val_loss: 0.0017 - val_root_mean_squared_error: 0.0408
Epoch 0041: val_root_mean_squared_error did not improve from 0.03956
Epoch 42/50 ===== - 2s 1ms/step - loss: 0.0257 - root_mean_squared_error: 0.1602 - val_loss: 0.0016 - val_root_mean_squared_error: 0.0398
Epoch 0042: val_root_mean_squared_error did not improve from 0.03956
Epoch 43/50 ===== - 2s 1ms/step - loss: 0.0260 - root_mean_squared_error: 0.1612 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0402
Epoch 0043: val_root_mean_squared_error did not improve from 0.03956
Epoch 44/50 ===== - 2s 1ms/step - loss: 0.0215 - root_mean_squared_error: 0.1466 - val_loss: 0.0010 - val_root_mean_squared_error: 0.0394
Epoch 0044: val_root_mean_squared_error improved from 0.03956 to 0.03947, saving model to model_0.h5
Epoch 45/50 ===== - 2s 1ms/step - loss: 0.0198 - root_mean_squared_error: 0.1407 - val_loss: 0.0016 - val_root_mean_squared_error: 0.0401
Epoch 0045: val_root_mean_squared_error improved from 0.03947 to 0.03935, saving model to model_0.h5
Epoch 46/50 ===== - 2s 1ms/step - loss: 0.0198 - root_mean_squared_error: 0.1407 - val_loss: 0.0016 - val_root_mean_squared_error: 0.0401
Epoch 0046: val_root_mean_squared_error improved from 0.03947 to 0.03935, saving model to model_0.h5
Epoch 47/50 ===== - 2s 1ms/step - loss: 0.0188 - root_mean_squared_error: 0.1370 - val_loss: 0.0018 - val_root_mean_squared_error: 0.0394
Epoch 0047: val_root_mean_squared_error did not improve from 0.03935
Epoch 48/50 ===== - 2s 1ms/step - loss: 0.0180 - root_mean_squared_error: 0.1342 - val_loss: 0.0015 - val_root_mean_squared_error: 0.0384
Epoch 0048: val_root_mean_squared_error did not improve from 0.03935
Epoch 49/50 ===== - 3s 2ms/step - loss: 0.0139 - root_mean_squared_error: 0.1180 - val_loss: 0.0011 - val_root_mean_squared_error: 0.0337
Epoch 0049: val_root_mean_squared_error improved from 0.03935 to 0.03944, saving model to model_0.h5
Epoch 50/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1183 - val_loss: 0.0010 - val_root_mean_squared_error: 0.0338
Epoch 0050: val_root_mean_squared_error improved from 0.03944 to 0.03937, saving model to model_0.h5
Epoch 51/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1138 - val_loss: 0.0011 - val_root_mean_squared_error: 0.0346
Epoch 0051: val_root_mean_squared_error improved from 0.03937 to 0.03465, saving model to model_0.h5
Epoch 52/50 ===== - 3s 2ms/step - loss: 0.0155 - root_mean_squared_error: 0.1246 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0351
Epoch 0052: val_root_mean_squared_error did not improve from 0.03465
Epoch 53/50 ===== - 3s 2ms/step - loss: 0.0140 - root_mean_squared_error: 0.1183 - val_loss: 0.0010 - val_root_mean_squared_error: 0.0338
Epoch 0053: val_root_mean_squared_error improved from 0.03465 to 0.03460, saving model to model_0.h5
Epoch 54/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1138 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0353
Epoch 0054: val_root_mean_squared_error improved from 0.03460 to 0.03465, saving model to model_0.h5
Epoch 55/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1138 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0353
Epoch 0055: val_root_mean_squared_error improved from 0.03465 to 0.03460, saving model to model_0.h5
Epoch 56/50 ===== - 3s 2ms/step - loss: 0.0137 - root_mean_squared_error: 0.1169 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0346
Epoch 0056: val_root_mean_squared_error improved from 0.03460 to 0.03426, saving model to model_0.h5
Epoch 57/50 ===== - 3s 2ms/step - loss: 0.0139 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0057: val_root_mean_squared_error did not improve from 0.03426
Epoch 58/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0058: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 59/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0059: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 60/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0060: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 61/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0061: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 62/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0062: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 63/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0063: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 64/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0064: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 65/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0065: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 66/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0066: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 67/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0067: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 68/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0068: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 69/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0069: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 70/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0070: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 71/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0071: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 72/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0072: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 73/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0073: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 74/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0074: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 75/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0075: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 76/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0076: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 77/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0077: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 78/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0078: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 79/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0079: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 80/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0080: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 81/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0081: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 82/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0082: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 83/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0083: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 84/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0084: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 85/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0085: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 86/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0086: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 87/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0087: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 88/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0088: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 89/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0089: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 90/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0090: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 91/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0091: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 92/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0092: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 93/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0093: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 94/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0094: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 95/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0095: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 96/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0096: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 97/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0097: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 98/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0098: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 99/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0099: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Epoch 100/50 ===== - 3s 2ms/step - loss: 0.0136 - root_mean_squared_error: 0.1134 - val_loss: 0.0012 - val_root_mean_squared_error: 0.0345
Epoch 0100: val_root_mean_squared_error improved from 0.03426 to 0.03426, saving model to model_0.h5
Out-of-Fold RMSE is: 0.03491139867002192

7. Test Model

Below I calculated root mean squared error on the test set.

In [21]: # Extract all test results
# This is the raw output of the model
test_pred = np.array(test_pred.reshape(-1))

# This is the result processed back to the number of contributors
test_pred_initial = np.round(math.e ** (test_pred * SCALE_COEFF))

In [22]: rmse = np.sqrt(mean_squared_error(test_observed, test_pred))
print(f'Testing root mean squared error: {rmse}')
Testing root mean squared error: 0.0345173300038445

Note

Since the label has been preprocessed, the root mean square of raw outputs does not tell us much about how the model is performing in real world situations, hence, I try to insert some other ways to help us visualize the performance better.

In the following cells, I will only work with post-processed predictions (predictions that has been transformed back to the number of contributors) a.k.a. the 'test_pred_initial' array.

In [23]: # This function tell us the percentage of predictions the model get right within a specified range
def test_accuracy(allowed_error, prediction, actual, verbose=True):
    corrects_percentage = np.mean(np.absolute(prediction - actual) <= allowed_error)
    if verbose:
        print(f'Around {round(corrects_percentage * 100)}% of predicted results are within {allowed_error} units
    return corrects_percentage

In [61]: # Some demonstration for the previous function
_ = test_accuracy(0, test_pred_initial, test_observed_initial)
_ = test_accuracy(1, test_pred_initial, test_observed_initial)

Around 43% of predicted results are within 0 units different from actual results.
Around 76% of predicted results are within 1 units different from actual results.

Conclude

So based on the previous demonstration, around 43% of the predictions are exactly the same as the actual observations; around 76% of the predictions are within 1 unit difference from the actual results. So given other features, the model can predict correctly the number of contributors for 76% of the cases, give or take 1 contributors.

Below is the graph plotting the relation between the allowed difference between predictions and observations and the percentage of the predictions are within that range.

In [25]: start = 0
stop = 30
step = 1
testing_range = np.arange(start, stop, step)
testing_accuracies = []
for e in testing_range:
    testing_accuracies.append(test_accuracy(e, test_pred_initial, test_observed_initial, verbose=False))

plt.plot(testing_range, testing_accuracies, label='Predictions')
plt.title('Accuracies base on Unit difference')
plt.xlabel('Unit difference')
plt.ylabel('Accuracy (%)')
plt.show()

Accuracies base on Unit difference

100
90
80
70
60
50
40
30
20
10
0
0 5 10 15 20 25 30
Unit difference

Note

It is worth noting that the dataset for this problem is very skewed, the majority of the datapoints has very low value label, the appearance of large label values are so insignificant that might cause the model to ignore all the large label observations or in other words, the model does not make any high prediction at all.

Hence, below I plot the predictions distribution and the relation between predictions and actual results to check for that. Just a smaller note, the points in the area between the 2 green lines in plot 2 is all the predictions that are within 3 unit difference from actual results.

In [26]: plt.subplots(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.hist(test_pred_initial, alpha=0.5, bins=40, color='red', label='Predictions')
plt.hist(test_observed_initial, alpha=0.5, bins=40, color='green', label='Observations')
plt.title('Predictions/Observations Distribution')
plt.grid(True)

plt.subplot(1, 2, 2)
plt.legend(loc='upper right')
plt.xlim(0, 20)
plt.xlabel('Contributors')
plt.ylabel('Count')

plt.subplot(1, 2, 2)
plt.scatter(test_pred_initial, test_observed_initial, alpha=0.5, color='blue')
unit_diff = 3
plt.plot([0, 35-unit_diff], [unit_diff, 35], alpha=0.8, color='green')
plt.plot([unit_diff, 35], [0, 35-unit_diff], alpha=0.8, color='green')
plt.title('Predictions/Observations Relation')
plt.grid(True)
plt.xlim(0, 35)
plt.ylim(0, 35)
plt.xlabel('Predicted value')
plt.ylabel('Observed value')
plt.show()

Predictions/Observations Distribution Predictions/Observations Relation

1200
1000
800
600
400
200
0
0 10 20 30 40 50 60 70
Contributors

35
30
25
20
15
10
5
0
0 5 10 15 20 25 30 35
Predicted value

Observed value

Conclude

From the plot, it is clear that the model does not just predict all small labels, and it is actually able to notice the patterns in some of the large labels (74y).

8. Conclusion

With only 5,800 datapoints and a shallow net structure, the model has been able to notice some patterns in large/small labels. Hence, to improve the model and solve this problem more thoroughly, here is my thoughts on what to do next:

Fetch more data from the API, preferably somewhere around 1 million datapoints so the model can have around 1~10 thousand exceptional (rare)-label data to learn from.

Extend API rate limitation to fetch more features from the API beside the basic info that I used.

Try out more complex functions than Logarithm to normalize the data distributions and values.

Make the neural net deeper and larger so the model can recognize more complex pattern.

Try more complex structure for the model, such as residual nets, or convolution nets, etc.

Try pretrained model.

Thank you for
```